

# Building a Robot Judge: Data Science for Decision-Making

## 3. Machine Learning Essentials

## Recap: Week 1 Activity

<https://padlet.com/eash44/53afqyu00cqsrrfc>

# Outline

## Essentials

Regression / Regularization

Activity on Causal Graphs

Binary Classification

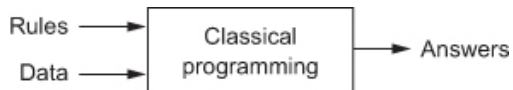
Applications

Appendix on Course Projects


# Learning Objectives

1. **Implement and evaluate machine learning pipelines.**
  - **Evaluate (find problems in) existing machine learning pipelines.**
  - **Design a pipeline to solve a given ML problem.**
  - **Implement some standard pipelines in Python.**
2. Implement and evaluate causal inference designs.
3. Understand how (not) to use data science tools (ML and CI) to support expert decision-making.

# What is machine learning?

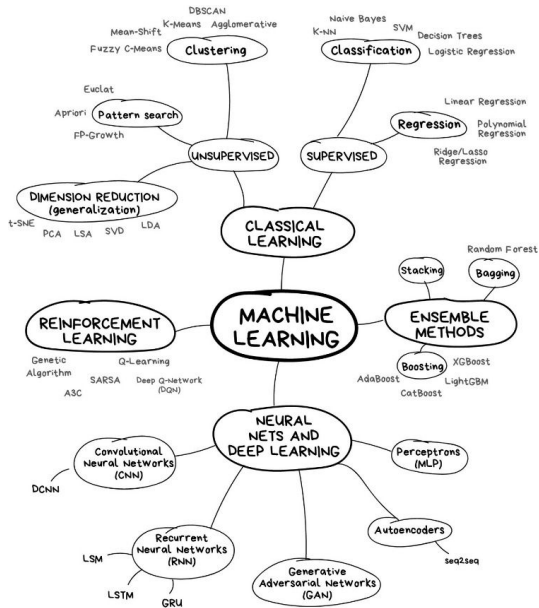


► In classical computer programming, humans input the rules and the data, and the computer provides answers.




► In machine learning, humans input the data and the answers, and the computer learns the rules.

# The Machine Learning Landscape



# What we will do today



# A Machine Learning Project, End-to-End

Aurelien Geron, *Hands-on machine learning with Scikit-Learn, Keras, & TensorFlow*, Chapter 2:

1. Look at the big picture.
2. Get the data.
3. Discover and visualize the data to gain insights.
4. Prepare the data for Machine Learning algorithms.
5. Select a model and train it.
6. Fine-tune your model.
7. Present your solution.
8. Launch, monitor, and maintain your system.



# Three Types of (Standard) Machine Learning Problems

Determined by the data type of the outcome variable (or label):

- ▶ **Binary classification**: two choices, normalized to zero and one.
  - ▶ e.g., guilty or innocent

# Three Types of (Standard) Machine Learning Problems

Determined by the data type of the outcome variable (or label):

- ▶ **Binary classification**: two choices, normalized to zero and one.
  - ▶ e.g., guilty or innocent
- ▶ **Regression**: a one-dimensional, continuous, real-valued outcome.
  - ▶ e.g., number of days of prison assigned

# Three Types of (Standard) Machine Learning Problems

Determined by the data type of the outcome variable (or label):

- ▶ **Binary classification:** two choices, normalized to zero and one.
  - ▶ e.g., guilty or innocent
- ▶ **Regression:** a one-dimensional, continuous, real-valued outcome.
  - ▶ e.g., number of days of prison assigned
- ▶ **Multinomial Classification:** Three or more discrete, un-ordered outcomes.
  - ▶ e.g., predict what judge is assigned to a case: Alito, Breyer, or Cardozo

## What type of ML Problem is this?

- ▶ Based on defendant characteristics and the facts of the case, predict which charges the prosecutor will bring:
  - ▶ third degree murder (manslaughter)
  - ▶ second degree murder (crime of passion)
  - ▶ first degree murder (premeditated)

**{binary classification, regression, or multinomial classification}**
- ▶ **write down your answer (30 secs)**

What do ML Algorithms do? Minimize a cost function

## What do ML Algorithms do? Minimize a cost function

- ▶ A typical cost function (or loss function) for regression problems is Mean Squared Error (MSE):

$$\text{MSE}(\theta) = \frac{1}{n_D} \sum_{i=1}^{n_D} (h(x_i; \theta) - y_i)^2$$

- ▶  $n_D$ , the number of rows/observations
- ▶  $x$ , the matrix of predictors, with row  $x_i$
- ▶  $y$ , the vector of outcomes, with item  $y_i$
- ▶  $h(x_i; \theta) = \hat{y}$  the model prediction (hypothesis)

## Loss functions, more generally

- ▶ The loss function  $L(\hat{\mathbf{y}}, \mathbf{y})$  assigns a score based on prediction and truth:
  - ▶ Should be bounded from below, with the minimum attained only for cases where the prediction is correct.
- ▶ The average loss for the test set is

$$\mathcal{L}(\theta) = \frac{1}{n_D} \sum_{i=1}^{n_D} L(h(\mathbf{x}_i; \theta), \mathbf{y}_i)$$

- ▶ The estimated parameter matrix  $\theta$  solves

$$\hat{\theta} = \arg \min_{\theta} \mathcal{L}(\theta)$$

↪ optimizes over parameter space; treats the data as constants.

# OLS Regression is Machine Learning

- ▶ Ordinary Least Squares Regression (OLS), also called simple linear regression, assumes the functional form  $h(x; \theta) = x_i' \theta$  and minimizes the mean squared error (MSE)

$$\min_{\hat{\theta}} \frac{1}{n_D} \sum_{i=1}^{n_D} (x_i' \hat{\theta} - y_i)^2$$



# OLS Regression is Machine Learning

- ▶ Ordinary Least Squares Regression (OLS), also called simple linear regression, assumes the functional form  $h(x; \theta) = x'_i \theta$  and minimizes the mean squared error (MSE)

$$\min_{\hat{\theta}} \frac{1}{n_D} \sum_{i=1}^{n_D} (x'_i \hat{\theta} - y_i)^2$$

- ▶ This minimand has a closed form solution

$$\hat{\theta} = (\mathbf{x}'\mathbf{x})^{-1} \mathbf{x}'\mathbf{y}$$

- ▶ most machine learning models do **not** have a closed form solution → use numerical optimization (gradient descent).

$$\text{MSE}(\theta) = \frac{1}{n_D} \sum_{i=1}^{n_D} (h(\theta; \mathbf{x}_i) - y_i)^2$$

- The partial derivative for feature  $j$  is

$$\frac{\partial \text{MSE}}{\partial \theta_j} = \frac{2}{n_D} \sum_{i=1}^{n_D} \underbrace{(h(\theta; \mathbf{x}_i) - y_i)}_{\text{error for this obs}} \underbrace{\frac{\partial h(\theta; \mathbf{x}_i)}{\partial \theta_j}}_{\text{how } \theta_j \text{ shifts } h(\cdot)}$$

- → estimates how changing  $\theta_j$  would reduce the error across the whole dataset.

$$\text{MSE}(\theta) = \frac{1}{n_D} \sum_{i=1}^{n_D} (h(\theta; \mathbf{x}_i) - y_i)^2$$

- ▶ The partial derivative for feature  $j$  is

$$\frac{\partial \text{MSE}}{\partial \theta_j} = \frac{2}{n_D} \sum_{i=1}^{n_D} \underbrace{(h(\theta; \mathbf{x}_i) - y_i)}_{\text{error for this obs}} \underbrace{\frac{\partial h(\theta; \mathbf{x}_i)}{\partial \theta_j}}_{\text{how } \theta_j \text{ shifts } h(\cdot)}$$

- ▶ → estimates how changing  $\theta_j$  would reduce the error across the whole dataset.
- ▶ The **gradient**  $\nabla$  gives the vector of these partial derivatives for all features:

$$\nabla_{\theta} \text{MSE} = \begin{bmatrix} \frac{\partial \text{MSE}}{\partial \theta_1} \\ \frac{\partial \text{MSE}}{\partial \theta_2} \\ \vdots \\ \frac{\partial \text{MSE}}{\partial \theta_j} \end{bmatrix}$$

$$\text{MSE}(\theta) = \frac{1}{n_D} \sum_{i=1}^{n_D} (h(\theta; \mathbf{x}_i) - y_i)^2$$

- ▶ The partial derivative for feature  $j$  is

$$\frac{\partial \text{MSE}}{\partial \theta_j} = \frac{2}{n_D} \sum_{i=1}^{n_D} \underbrace{(h(\theta; \mathbf{x}_i) - y_i)}_{\text{error for this obs}} \underbrace{\frac{\partial h(\theta; \mathbf{x}_i)}{\partial \theta_j}}_{\text{how } \theta_j \text{ shifts } h(\cdot)}$$

- ▶ → estimates how changing  $\theta_j$  would reduce the error across the whole dataset.

- ▶ The **gradient**  $\nabla$  gives the vector of these partial derivatives for all features:

$$\nabla_{\theta} \text{MSE} = \begin{bmatrix} \frac{\partial \text{MSE}}{\partial \theta_1} \\ \frac{\partial \text{MSE}}{\partial \theta_2} \\ \vdots \\ \frac{\partial \text{MSE}}{\partial \theta_j} \end{bmatrix}$$

- ▶ **Gradient descent** nudges  $\theta$  against the gradient (the direction that reduces MSE):

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \text{MSE}$$

- ▶  $\eta$  = learning rate
- ▶ keep nudging until convergence.

$$\text{MSE}(\theta) = \frac{1}{n_D} \sum_{i=1}^{n_D} (h(\theta; \mathbf{x}_i) - y_i)^2$$

- ▶ The partial derivative for feature  $j$  is

$$\frac{\partial \text{MSE}}{\partial \theta_j} = \frac{2}{n_D} \sum_{i=1}^{n_D} \underbrace{(h(\theta; \mathbf{x}_i) - y_i)}_{\text{error for this obs}} \underbrace{\frac{\partial h(\theta; \mathbf{x}_i)}{\partial \theta_j}}_{\text{how } \theta_j \text{ shifts } h(\cdot)}$$

- ▶ → estimates how changing  $\theta_j$  would reduce the error across the whole dataset.

- ▶ The **gradient**  $\nabla$  gives the vector of these partial derivatives for all features:

$$\nabla_{\theta} \text{MSE} = \begin{bmatrix} \frac{\partial \text{MSE}}{\partial \theta_1} \\ \frac{\partial \text{MSE}}{\partial \theta_2} \\ \vdots \\ \frac{\partial \text{MSE}}{\partial \theta_j} \end{bmatrix}$$

- ▶ **Gradient descent** nudges  $\theta$  against the gradient (the direction that reduces MSE):

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \text{MSE}$$

- ▶  $\eta$  = learning rate

- ▶ keep nudging until convergence.

- ▶ **Stochastic gradient descent (SGD)**: Compute gradient for single random instance (rather than whole dataset) at each iteration. Much faster, still works.

# Data Prep for Machine Learning

- ▶ Data Pre-Processing: See Geron Chapter 2 for pandas and sklearn syntax:
  - ▶ imputing missing values.
  - ▶ feature scaling (often helpful/necessary for ML models to work well)
    - ▶ if predictors are sparse (e.g. bag-of-words), use `StandardScaler(with_mean=False)`.
  - ▶ encoding categorical variables.
  - ▶ **Best practice: reproducible data pipeline.**

# Data Prep for Machine Learning

- ▶ Data Pre-Processing: See Geron Chapter 2 for pandas and sklearn syntax:
  - ▶ imputing missing values.
  - ▶ feature scaling (often helpful/necessary for ML models to work well)
    - ▶ if predictors are sparse (e.g. bag-of-words), use `StandardScaler(with_mean=False)`.
  - ▶ encoding categorical variables.
  - ▶ **Best practice: reproducible data pipeline.**
- ▶ Train/Test Split:
  - ▶ ML models can achieve arbitrarily high accuracy in-sample, so performance should be evaluated out-of-sample.

# Data Prep for Machine Learning

- ▶ Data Pre-Processing: See Geron Chapter 2 for pandas and sklearn syntax:
  - ▶ imputing missing values.
  - ▶ feature scaling (often helpful/necessary for ML models to work well)
    - ▶ if predictors are sparse (e.g. bag-of-words), use `StandardScaler(with_mean=False)`.
  - ▶ encoding categorical variables.
  - ▶ **Best practice: reproducible data pipeline.**
- ▶ Train/Test Split:
  - ▶ ML models can achieve arbitrarily high accuracy in-sample, so performance should be evaluated out-of-sample.
  - ▶ standard approach: randomly sample 80% training dataset to learn parameters, form predictions in 20% testing dataset for evaluating performance.



# Use Cross-Validation During Model Training

- ▶ Within the training set:
  - ▶ Use cross-validation with grid search to get model performance metrics across subsets of data using different hyperparameter specs.
  - ▶ Find the best hyperparameters for out-of-fold prediction in the training set.
- ▶ Then evaluate model performance in the test set using these hyperparameters.

# Use Cross-Validation During Model Training

- ▶ Within the training set:
  - ▶ Use cross-validation with grid search to get model performance metrics across subsets of data using different hyperparameter specs.
  - ▶ Find the best hyperparameters for out-of-fold prediction in the training set.
- ▶ Then evaluate model performance in the test set using these hyperparameters.
- ▶ Cross-validation is less common in deep learning, where training multiple models is too computationally expensive.
  - ▶ instead, use dropout and early stopping (week 7).

# Model Evaluation in Test Set

Evaluating a “good” model is context-dependent. Here are some basics.

Regression:

- ▶ mean squared error (MSE)
- ▶ R-squared (same ranking as MSE, but units are more interpretable)
- ▶ mean absolute error (MAE,  $\sum |\hat{y}(\theta) - y|$ ) is less sensitive to outliers.

# Model Evaluation in Test Set

Evaluating a “good” model is context-dependent. Here are some basics.

Regression:

- ▶ mean squared error (MSE)
- ▶ R-squared (same ranking as MSE, but units are more interpretable)
- ▶ mean absolute error (MAE,  $\sum |\hat{y}(\theta) - y|$ ) is less sensitive to outliers.

Classification:

- ▶ more complicated, but accuracy is a good baseline:  
accuracy = ( $\#$  correct test-set predictions) / ( $\#$  of test-set observations)

# Model Evaluation in Test Set

Evaluating a “good” model is context-dependent. Here are some basics.

Regression:

- ▶ mean squared error (MSE)
- ▶ R-squared (same ranking as MSE, but units are more interpretable)
- ▶ mean absolute error (MAE,  $\sum |\hat{y}(\theta) - y|$ ) is less sensitive to outliers.

Classification:

- ▶ more complicated, but accuracy is a good baseline:  
accuracy = ( $\#$  correct test-set predictions) / ( $\#$  of test-set observations)
- ▶ What if one of the outcomes is over-represented – e.g., 19 out of 20? Then I can guess the modal class and get 95% accuracy.
  - ▶ Some alternative classifier metrics designed to address class imbalance (more below and in week 5).

# Outline

Essentials

Regression / Regularization

Activity on Causal Graphs

Binary Classification

Applications

Appendix on Course Projects

## Regression models $\leftrightarrow$ Continuous outcome

- ▶ If the outcome is continuous (e.g.,  $Y$  = tax revenues collected, or criminal sentence imposed in months of prison):
- ▶ Need a regression model. Problems with OLS:
  - ▶ tends to over-fit training data.
  - ▶ cannot handle multicollinearity.




Figure 4-14. High-degree Polynomial Regression

## Regression models $\leftrightarrow$ Continuous outcome

- ▶ If the outcome is continuous (e.g.,  $Y$  = tax revenues collected, or criminal sentence imposed in months of prison):

- ▶ Need a regression model. Problems with OLS:

- ▶ tends to over-fit training data.
- ▶ cannot handle multicollinearity.




Figure 4-14. High-degree Polynomial Regression

- ▶ Machine learning models are evaluated by the fit in held-out data (the test set)
  - ▶ “Regularization” refers to ML model training methods designed to reduce/prevent over-fitting of the training set
  - ▶ (and hopefully better fit in the test set).



## Regularization

- ▶ Minimizing the loss  $L$  directly usually results in over-fitting. It is standard to add regularization:

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n_D} \sum_{i=1}^{n_D} L(h(\mathbf{x}_i; \theta), \mathbf{y}_i) + \lambda R(\theta)$$

- ▶  $R(\theta)$  is a “regularization function” or “regularizer”, designed to reduce over-fitting.
  - ▶  $\lambda$  is a hyperparameter where higher values increase regularization.
-

## Regularization

- ▶ Minimizing the loss  $L$  directly usually results in over-fitting. It is standard to add regularization:

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n_D} \sum_{i=1}^{n_D} L(h(\mathbf{x}_i; \theta), \mathbf{y}_i) + \lambda R(\theta)$$

- ▶  $R(\theta)$  is a “regularization function” or “regularizer”, designed to reduce over-fitting.
- ▶  $\lambda$  is a hyperparameter where higher values increase regularization.

---

“Ridge” and “Lasso” penalize larger coefficients, shrinking them toward zero:

- ▶ Ridge (or L2) penalty:

$$R_2 = \|\theta\|_2^2 = \sum_{j=1}^{n_x} (\theta_j)^2$$

- ▶ also helps select between collinear predictors.

- ▶ Lasso (or L1) penalty:


$$R_1 = \|\theta\|_1 = \sum_{j=1}^{n_x} |\theta_j|$$

- ▶ also performs feature selection and outputs a sparse model.

# Does lasso pick the “true” model?

Lasso prediction of house prices with 150 variables – which variables are “selected” (non-zero coefficients) by lasso, in ten models trained on separate data subsamples (Mullainathan and Spiess 2017):


Selected Coefficients (Nonzero Estimates) across Ten LASSO Regressions



# Does lasso pick the “true” model?

Lasso prediction of house prices with 150 variables – which variables are “selected” (non-zero coefficients) by lasso, in ten models trained on separate data subsamples (Mullainathan and Spiess 2017):

Selected Coefficients (Nonzero Estimates) across Ten LASSO Regressions



- ▶ The set of lasso-selected variables changes across folds in the data
- ▶ → Lasso does not pick the “correct” predictors.
  - ▶ It just learns the correct  $\hat{h}(X)$
  - ▶ when predictors are correlated with each other, they are substitutable.

## Elastic Net = Lasso + Ridge

The Elastic Net cost function is:

$$\begin{aligned} L(\theta) &= \text{MSE}(\theta) + \lambda_1 R_1 + \lambda_2 R_2 \\ &= \text{MSE}(\theta) + \lambda_1 \sum_{j=1}^{n_x} |\theta_j| + \lambda_2 \sum_{j=1}^{n_x} (\theta_j)^2 \end{aligned}$$

- ▶  $\lambda_1, \lambda_2$  = strength of L1 (Lasso) penalty and L2 (Ridge) penalty, respectively.
-

## Elastic Net = Lasso + Ridge

The Elastic Net cost function is:

$$\begin{aligned} L(\theta) &= \text{MSE}(\theta) + \lambda_1 R_1 + \lambda_2 R_2 \\ &= \text{MSE}(\theta) + \lambda_1 \sum_{j=1}^{n_x} |\theta_j| + \lambda_2 \sum_{j=1}^{n_x} (\theta_j)^2 \end{aligned}$$

- ▶  $\lambda_1, \lambda_2$  = strength of L1 (Lasso) penalty and L2 (Ridge) penalty, respectively.

---

In scikit-learn, e-net penalties are parametrized as “alpha” = total penalty, and “l1\_ratio” = proportion of penalty to L1.

```
from sklearn.linear_model import ElasticNet
enet = ElasticNet(alpha=2.0, l1_ratio = .75) # L1 = 1.5, L2 = 0.5
enet.fit(X,y)
```

# Outline

Essentials

Regression / Regularization


Activity on Causal Graphs

Binary Classification

Applications


Appendix on Course Projects

## Causal Graphs: Review





## Causal Graph Example: Pollution of a River



## Activity: Practice with Causal Graphs

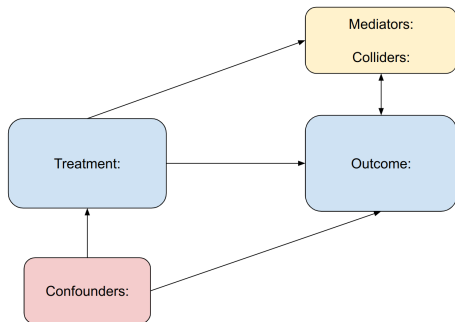
- ▶ Think of an example causal inference question:
  - ▶ a research question from your field
  - ▶ a policy you are interested in
  - ▶ a mystery you are fascinated by

## Activity: Practice with Causal Graphs

- ▶ Think of an example causal inference question:
  - ▶ a research question from your field
  - ▶ a policy you are interested in
  - ▶ a mystery you are fascinated by

Causal graph template

(also at <http://bit.ly/BRJ-W5A2a>):



- ▶ Draw this template on a piece of paper (recommended), or save a copy of the template
  - ▶ fill it in on paper or electronically
  - ▶ after the break, we will share in groups of 2-3

# Outline

Essentials

Regression / Regularization

Activity on Causal Graphs

**Binary Classification**

Applications

Appendix on Course Projects

# Binary Outcome $\leftrightarrow$ Binary Classification

- ▶ Binary classifiers try to match a boolean outcome  $y \in \{0, 1\}$ .
  - ▶ The standard approach is to apply a transformation (e.g. sigmoid/logit) to normalize  $\hat{y} \in [0, 1]$ .
  - ▶ Prediction rule is 0 for  $\hat{y} < .5$  and 1 otherwise.


## Binary Outcome $\leftrightarrow$ Binary Classification

- ▶ Binary classifiers try to match a boolean outcome  $y \in \{0, 1\}$ .
  - ▶ The standard approach is to apply a transformation (e.g. sigmoid/logit) to normalize  $\hat{y} \in [0, 1]$ .
  - ▶ Prediction rule is 0 for  $\hat{y} < .5$  and 1 otherwise.
- ▶ The binary cross-entropy (or log loss) is:

$$L(\theta) = \underbrace{-\frac{1}{n_D}}_{\text{negative}} \sum_{i=1}^{n_D} \left[ \underbrace{y_i}_{y_i=1} \underbrace{\log(\hat{y}_i)}_{\log \text{ prob}_{y_i=1}} + \underbrace{(1-y_i)}_{y_i=0} \underbrace{\log(1-\hat{y}_i)}_{\log \text{ prob}_{y_i=0}} \right]$$


- In **logistic regression** we use a sigmoid transformation:

$$\hat{y} = \text{sigmoid}(\mathbf{x} \cdot \theta) = \frac{1}{1 + \exp(-\mathbf{x} \cdot \theta)}$$



- In **logistic regression** we use a sigmoid transformation:

$$\hat{y} = \text{sigmoid}(\mathbf{x} \cdot \theta) = \frac{1}{1 + \exp(-\mathbf{x} \cdot \theta)}$$



- Plugging into the binary-cross entropy loss gives the logistic regression cost objective:


$$\min_{\theta} \sum_{i=1}^{n_D} -y_i \log(\text{sigmoid}(\mathbf{x}_i \cdot \theta)) - [1 - y_i] \log(1 - \text{sigmoid}(\mathbf{x}_i \cdot \theta))$$

- does not have a closed form solution, but it is convex (guaranteeing that gradient descent will find the global minimum).



- ▶ In **logistic regression** we use a sigmoid transformation:

$$\hat{y} = \text{sigmoid}(\mathbf{x} \cdot \theta) = \frac{1}{1 + \exp(-\mathbf{x} \cdot \theta)}$$



- ▶ Plugging into the binary-cross entropy loss gives the logistic regression cost objective:


$$\min_{\theta} \sum_{i=1}^{n_D} -y_i \log(\text{sigmoid}(\mathbf{x}_i \cdot \theta)) - [1 - y_i] \log(1 - \text{sigmoid}(\mathbf{x}_i \cdot \theta))$$

- ▶ does not have a closed form solution, but it is convex (guaranteeing that gradient descent will find the global minimum).
- ▶ The gradient for one data point is

$$\frac{\partial L(\theta)}{\partial \theta_j} = \underbrace{(\text{sigmoid}(\mathbf{x}_i \cdot \theta) - y_i)}_{\text{error for obs } i} \underbrace{x_i^j}_{\text{input } j}$$

- ▶ In **logistic regression** we use a sigmoid transformation:

$$\hat{y} = \text{sigmoid}(\mathbf{x} \cdot \theta) = \frac{1}{1 + \exp(-\mathbf{x} \cdot \theta)}$$



- ▶ Plugging into the binary-cross entropy loss gives the logistic regression cost objective:

$$\min_{\theta} \sum_{i=1}^{n_D} -y_i \log(\text{sigmoid}(\mathbf{x}_i \cdot \theta)) - [1 - y_i] \log(1 - \text{sigmoid}(\mathbf{x}_i \cdot \theta))$$

- ▶ does not have a closed form solution, but it is convex (guaranteeing that gradient descent will find the global minimum).
- ▶ The gradient for one data point is

$$\frac{\partial L(\theta)}{\partial \theta_j} = \underbrace{(\text{sigmoid}(\mathbf{x}_i \cdot \theta) - y_i)}_{\text{error for obs } i} \underbrace{x_i^j}_{\text{input } j}$$

- ▶ Like linear regression, logistic regression can be regularized with L1 or L2 penalties.

```
from sklearn.linear_model import LogisticRegression
logit = LogisticRegression(penalty='l2', C = 2.0) # lambda = 1/2
logit.fit(X,y)
```

A **Confusion Matrix** is a nice way to visualize classifier performance:

		Predicted Class	
		Negative	Positive
True Class	Negative	<b># True Negatives</b>	# False Positives
	Positive	# False Negatives	<b># True Positives</b>

- ▶ Cell values give counts in the test set.

A **Confusion Matrix** is a nice way to visualize classifier performance:

		Predicted Class	
		Negative	Positive
True Class	Negative	<b># True Negatives</b>	# False Positives
	Positive	# False Negatives	<b># True Positives</b>

- ▶ Cell values give counts in the test set.

$$\text{Accuracy} = \frac{\text{True Positives} + \text{True Negatives}}{\text{True Positives} + \text{False Positives} + \text{False Negatives} + \text{True Negatives}}$$

A **Confusion Matrix** is a nice way to visualize classifier performance:

		Predicted Class	
		Negative	Positive
True Class	Negative	<b># True Negatives</b>	# False Positives
	Positive	# False Negatives	<b># True Positives</b>

- ▶ Cell values give counts in the test set.

$$\text{Accuracy} = \frac{\text{True Positives} + \text{True Negatives}}{\text{True Positives} + \text{False Positives} + \text{False Negatives} + \text{True Negatives}}$$

$$\text{Precision (for positive class)} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

- ▶ Precision decreases with false positives. “When I guess this outcome, I tend to guesses correctly.”

A **Confusion Matrix** is a nice way to visualize classifier performance:

		Predicted Class	
		Negative	Positive
True Class	Negative	<b># True Negatives</b>	# False Positives
	Positive	# False Negatives	<b># True Positives</b>

- ▶ Cell values give counts in the test set.

$$\text{Accuracy} = \frac{\text{True Positives} + \text{True Negatives}}{\text{True Positives} + \text{False Positives} + \text{False Negatives} + \text{True Negatives}}$$

$$\text{Precision (for positive class)} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

- ▶ Precision decreases with false positives. “When I guess this outcome, I tend to guesses correctly.”

$$\text{Recall (for positive class)} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

- ▶ Recall decreases with false negatives. “When this outcome occurs, I don’t miss it.”

# Outline

Essentials

Regression / Regularization

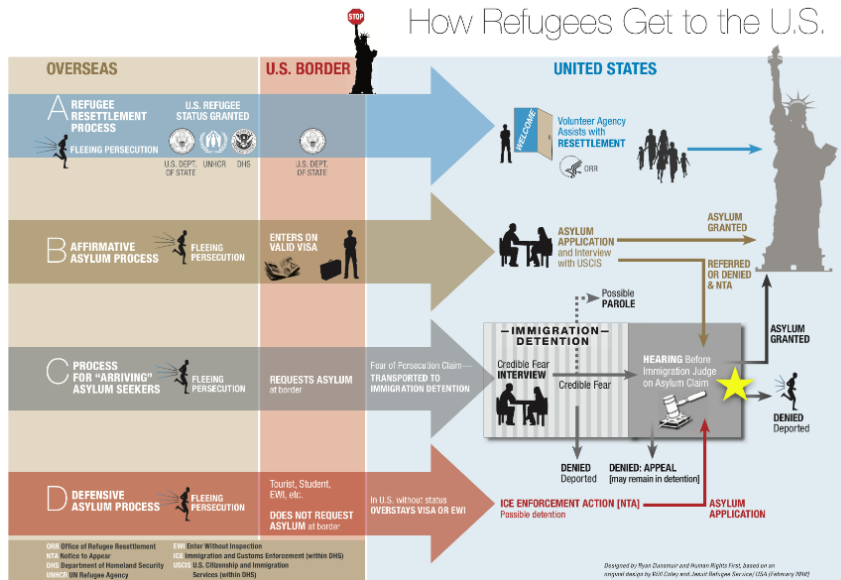
Activity on Causal Graphs

Binary Classification

Applications

Appendix on Course Projects

# Asylum in the U.S.





## Dunn, Sagun, Sirin, and Chen (2017): Asylum Courts

- ▶ Data:
  - ▶ universe of asylum court cases, 1981-2013
  - ▶ 492,903 decisions, 336 courts, 441 judges

## Dunn, Sagun, Sirin, and Chen (2017): Asylum Courts

- ▶ Data:
  - ▶ universe of asylum court cases, 1981-2013
  - ▶ 492,903 decisions, 336 courts, 441 judges
- ▶ High stakes: denial of asylum results in deportation.
- ▶ Average grant rate: 35%.
- ▶ What type of ML problem is this?

# Predicting U.S. Asylum Court Decisions

		Predicted	
		Denied	Granted
True	Denied	195,223	65,798
	Granted	73,269	104,406

**Accuracy = 68.3%, F1 = 0.60**

- ▶ Prediction App: <https://floating-lake-11821.herokuapp.com/>
  - ▶ predictions made using logistic regression with L2 regularization, penalty selected by cross-validation grid search.

## Judge Identity is Most Predictive Factor

Model	Accuracy	ROC AUC
Judge ID	0.71	0.74
Judge ID & Nationality	0.76	0.82
Judge ID & Opening Date	0.73	0.77
Judge ID & Nationality & Opening Date	0.78	0.84
Full model at case completion	0.82	0.88

- ▶ Predictions from random forest classifier, with parameters selected by cross-validated grid search.
  - ▶ Training/test split 482K/120K.

## Judge Variation in Predictability

- ▶ Some judges are highly predictable, always granting or rejecting.
  - ▶ suggests they use heuristics or stereotypes rather than considering cases carefully.

## Judge Variation in Predictability

- ▶ Some judges are highly predictable, always granting or rejecting.
  - ▶ suggests they use heuristics or stereotypes rather than considering cases carefully.
- ▶ There is significant variation in predictability by judge, conditional on grant rate.
  - ▶ suggests disagreement about circumstances contributing to asylum decision.

# What type of ML Problem is this? (Bonica 2018)

[http://bit.ly/BRJ\\_bonica](http://bit.ly/BRJ_bonica)

**Abstract:** *This article develops a generalized supervised learning methodology for inferring roll-call scores from campaign contribution data. Rather than use unsupervised methods to recover a latent dimension that best explains patterns in giving, donation patterns are instead mapped onto a target measure of legislative voting behavior. Supervised models significantly outperform alternative measures of ideology in predicting legislative voting behavior. Fundraising prior to entering office provides a highly informative signal about future voting behavior. Impressively, forecasts based on fundraising as a nonincumbent predict future voting behavior as accurately as in-sample forecasts based on votes cast during a legislator's first 2 years in Congress. The combined results demonstrate campaign contributions are powerful predictors of roll-call voting behavior and resolve an ongoing debate as to whether contribution data successfully distinguish between members of the same party.*

- **read and make a first decision (3 minutes); discuss in groups from earlier (2 minutes), then show of hands.**

## Vaccine Allocation, Age, and Race (rest of class)

- ▶ In groups: Discuss Noah Smith's article, "Vaccine Allocation, Age, and Race"
  - ▶ Summarize the main points, and relate them to our discussions in class.
  - ▶ Think of another policy/decision with similar issues, explain.
- ▶ Task:
  - ▶ Discuss your ideas in your group
  - ▶ write them down in a shareable doc
  - ▶ post a link in the padlet: <https://padlet.com/eash44/58p73xm4njmqcg7e>.



# Outline

Essentials

Regression / Regularization

Activity on Causal Graphs

Binary Classification

Applications

Appendix on Course Projects

# Course Project Logistics

<https://eash.cc/brj22-proj>

- ▶ If you are signed up for the credits, the focus of your work in this course should be on the project.
  - ▶ Can be done individually or in small groups (up to 4 students).
  - ▶ Do an original analysis using methods learned in the course, and write a paper about it.

# Course Project Logistics

<https://eash.cc/brj22-proj>

- ▶ If you are signed up for the credits, the focus of your work in this course should be on the project.
  - ▶ Can be done individually or in small groups (up to 4 students).
  - ▶ Do an original analysis using methods learned in the course, and write a paper about it.
- ▶ Deliverables (dates at project page)
  - ▶ description of topic (10% completion grade)
  - ▶ proposal/outline (10% of grade)
  - ▶ Presentations/Posters (10% of grade).
  - ▶ Rough draft with data/methods/results (20% of grade)
  - ▶ Final draft with full text and replication package (50% of grade)

## Previous Year's Projects (1)

## Previous Year's Projects (1)

- ▶ One of the groups began building a legal research application for Swiss lawyers:
  - ▶ see <https://deepjudge.ai/>
  - ▶ feature-rich legal search engine, won some VC funding and now part of ETH AI Center

## Previous Year's Projects (1)

- ▶ One of the groups began building a legal research application for Swiss lawyers:
  - ▶ see <https://deepjudge.ai/>
  - ▶ feature-rich legal search engine, won some VC funding and now part of ETH AI Center
- ▶ Another group partnered with a local company to build out environmental-regulation analytics
  - ▶ won an Innosuisse grant.

## Previous Year's Projects (2)

### Five projects have been published:

1. “Legal language modeling with transformers” (Lazar Peric, Stefan Mijic, Dominik Stammbach, Elliott Ash), *Proceedings of ASAIL* (2020).
2. “Entropy in Legal Language” (Roland Friedrich, Mauro Luzzatto, Elliott Ash), *NLLP @ KDD* (2020).
3. “Towards Automated Anamnesis Summarization: BERT-based Models for Symptom Extraction” (Anton Schäfer, Nils Blach, Oliver Rausch, Maximilian Warm, Nils Krüger), *Machine Learning for Health at NeurIPS* (2020).
4. “Kwame: A Bilingual AI Teaching Assistant for Online SuaCode Courses” (George Boateng), *International Conference on AI in Education* (2021)
5. “MemSum: Extractive Summarization of Long Documents using Multi-Step Episodic Markov Decision Processes” (Nianlong Gu, Elliott Ash, Richard Hahnloser), forthcoming ACL Main Conference (2022)

## Previous Year's Projects (3)

A number of other projects that are likely to get published, e.g.:

1. partisan tweet generator that responds in the style of a Republican or Democrat.
2. analysis of bias towards immigrants in the early 1900s using old newspapers.
3. causal analysis using deep instrumental variables of what arguments in judicial opinions increase citations
4. partisan question answering system that answers questions with a partisan slant.
5. an audio/text analysis of central bank speeches and inflation beliefs.
6. system for predicting judicial decisions based on the submitted briefs
7. analysis of sentiment in Tweet images, relate to local economic outcomes
8. automated determination of trademarkability
9. conditioned text generation with quality filtering
10. improvements to relatio narrative extraction engine



# Project Topics and First Steps

- ▶ Picking a topic:
  - ▶ You are welcome to come up with your own topic. We will provide feedback on that.
  - ▶ We have a list of suggested topics with project advisors.
  - ▶ I can also provide advice about which of these topics is a good fit based on team interests and skills.
- ▶ First steps:
  - ▶ once you have formed a group, send to Afra a list of team members with their resumes, research experience, and interests.
    - ▶ if you are interested in one or more of the suggested topics, include that in the email
  - ▶ we will then match project advisors and set up meeting

## Questions / comments?

- ▶ As suggested, we will set up a meet-and-greet for those doing projects.